

Lead Triage — how it works

A short note on what I built, how it works, and why I made some of the calls I made.

What this thing does

You drop a CSV of inbound leads inside — contact details, company info, and the free-text notes from previous chats. The system gives you back a ranked shortlist with three buckets: **contact now**, **nurture**, and **disqualify**. It runs as a Next.js app. You put the CSV in the upload box, you get a scored table back. The sample export (520 rows) runs in about 200ms on the server. No need to look at anything by hand.

How I built it

It is a single-page Next.js 16 + TypeScript app on Vercel. The flow is simple: CSV → POST /api/triage → JSON → table. No database, no login, no external API calls. Everything runs on the free Vercel tier. The scoring logic sits in `src/lib/leads/` as plain functions inside four files: `cleaners.ts` handles date, employee, budget and ID normalisation; `signals.ts` holds the phrase-banks and signal extractor; `score.ts` has the decision rules; `triage.ts` is the orchestrator that ties everything together. Every weight is a named constant, so you can check and tune it yourself.

How you use it: 1) Click *Use sample data* or upload your CSV. 2) The server parses, cleans, scores and ranks every lead. 3) The UI shows the ranked table — contact now on top, sorted by score. 4) Filter by status, source, or just search by text. 5) Click any row to see the full signal breakdown and the original notes. 6) Click Export to download the triaged list as a CSV, with the top signals per lead.

What counts as worth contacting now

A lead is **contact now** only if it has a clear **urgency signal** in the notes *and* a score of 65 or above. Urgency signals are phrases like *budget approved*, *decision this month*, *ready to pilot*, *I make the call here*, *keen to move fast*. The score is just a sum of weighted signals:

Signal kind	Weight	Examples
Urgency	+18 to +28	budget approved; decision this month; ready to pilot; I make the call
Pain / need	+8 to +14	eating our week; by hand; want it automated end to end
Budget	+5 to +40	parsed from \$8,000/mo, 5k-7k, 8k, 0, TBD
Authority (title)	-25 to +30	CEO/Founder +30 · VP/Head of +20 · Student/Recruiter -25
Company size	0 to +25	1 = 0 · 2-10 = +10 · 11-50 = +18 · 50+ = +25
Source	-40 to +22	referral +22 · linkedin +14 · webform +8 · junk -40
Disqualifier phrase	-100	looking for a role; attaching CV; VC here; SMM panel; WON \$1M

On the sample export of 520 leads: **100 contact now** (19%), **245 nurture** (47%), **175 disqualify** (34%). Average score 57.4, median 64.

Some of the calls I made

- The CSV format from the brief is what I treated as the standard input. Any future export with the same columns will just work.
- If somebody says comparing options but there is no urgency signal, I put them in nurture, not contact now. It could go either way, but I chose the careful side.
- VC here — wanting to intro portfolio companies — disqualifies. They are not the buyer.
- Date 04-06-2024 I read as D-M-Y, because I checked it against 19-06-2024 and 26-06-2024 rows in the export.
- For budget range like 5k-7k I use the lower bound (\$5k). Just to be safe.

Trade-offs I am okay with

- **Rules over LLM.** Same CSV always gives the same output, no API key, and every signal shows in the UI so you can see why. The flip side: it will not catch new phrasings — the keyword lists will need small tuning from time to time.
- **Contact-now is tight on purpose.** Budget and title alone is not enough — too many comparing-options leads have those. Asking for an urgency phrase keeps the list at about 19% of inbound, which a rep can actually work through in a week.
- **Disqualify is not permanent.** Leads stay in the export, they just get flagged. Easy to override.
- **Every signal is shown in the UI,** not just the score. So the rep can audit and override. The score is a starting point, not a final verdict.