

Lead Triage — System Documentation

A short note on what this is, how it works, and why it makes the calls it makes. Built for the Cohort 3 assessment.

Overview

The system takes a CSV export of inbound leads (contact details, company info, free-text notes from previous conversations) and returns a ranked shortlist with three buckets: **contact now**, **nurture**, **disqualify**. It runs as a Next.js app — drop the CSV in the upload zone, get a scored table back. The sample export (520 rows) processes in ~200ms server-side. No manual review required.

Architecture & how the workflow runs

Single-page Next.js 16 + TypeScript app on Vercel. Flow: CSV → POST /api/triage → JSON → table. No database, no auth, no external API calls — the whole thing runs on the free Vercel tier. The scoring logic lives in `src/lib/leads/` as pure functions split across four files: `cleaners.ts` (date / employee / budget / ID normalisation), `signals.ts` (the phrase-banks + signal extractor), `score.ts` (decision rules), `triage.ts` (orchestrator). Every weight is a named constant you can audit and tune.

Workflow: 1) Click *Use sample data* or upload a CSV. 2) Server parses, cleans, scores, ranks every lead. 3) UI shows ranked table — contact now at the top, sorted by score. 4) Filter by status / source / free-text search. 5) Click any row to see the full signal breakdown and original notes. 6) Click Export to download the triaged list as CSV with the top signals per lead.

Qualification logic — what counts as "worth contacting"

A lead is **contact now** only if it has an explicit **urgency signal** in the notes *and* a score ≥ 65 . Urgency signals are phrases like *budget approved*, *decision this month*, *ready to pilot*, *I make the call here*, *keen to move fast*. The score is a weighted sum of discrete signals:

Signal kind	Weight	Examples
Urgency	+18 to +28	budget approved; decision this month; ready to pilot; I make the call
Pain / need	+8 to +14	eating our week; by hand; want it automated end to end
Budget	+5 to +40	parsed from \$8,000/mo, 5k-7k, 8k, 0, TBD
Authority (title)	-25 to +30	CEO/Founder +30 · VP/Head of +20 · Student/Recruiter -25
Company size	0 to +25	$1 = 0 \cdot 2\text{-}10 = +10 \cdot 11\text{-}50 = +18 \cdot 50+ = +25$
Source	-40 to +22	referral +22 · linkedin +14 · webform +8 · junk -40
Disqualifier phrase	-100	looking for a role; attaching CV; VC here; SMM panel; WON \$1M

On the sample export of 520 leads: **100 contact now** (19%), **245 nurture** (47%), **175 disqualify** (34%). Average score 57.4, median 64.

Key assumptions

- The CSV format from the brief is the canonical input — future exports with the same columns just work.
- "Comparing options" + no urgency = nurture, not contact now. Could go either way; picked the conservative side.
- "VC here — wanting to intro portfolio companies" disqualifies. They're not the buyer.
- Date 04-06-2024 resolved as D-M-Y based on sampling the export (verified against 19-06-2024 and 26-06-2024 rows).
- Budget range "5k-7k" → use lower bound (\$5k). Conservative.

Design decisions & trade-offs

- **Rule-based over LLM.** Deterministic (same CSV → same output), no API key, every signal is explainable in the UI. Trade-off: won't catch novel phrasings — the keyword lists need occasional tuning.
- **Contact-now threshold is deliberately tight.** Budget + title alone doesn't qualify — too many "comparing options" leads have those. Requiring an urgency phrase keeps the list at ~19% of inbound, which a rep can actually work through in a week.
- **Disqualify is non-destructive.** Leads stay in the export, just flagged. Easy to override.
- **Every signal is surfaced in the UI**, not just the score. Lets the rep audit and override. The score is a starting point, not a verdict.